

CS7641 ML Practice Quiz

Module RL 1: Markov Decision Processes

Created by Kyle Nakamura with the help of ChatGPT

Question 1

What is a Markov Decision Process (MDP) primarily used for in reinforcement learning?

- A. To model deterministic environments where the outcome of each action is predictable.
- B. To represent decision-making scenarios that involve uncertainties and probabilistic outcomes.
- C. To optimize linear transformation algorithms for feature selection.
- D. To directly learn the policy which maps states to actions without using value functions.
- E. To solve problems where the transition and reward functions are known in advance.

Question 2

In the context of MDPs, what does the Markovian property imply about the transition function?

- A. It depends on the entire history of past states and actions.
- B. It only depends on the current state and not on any past states.
- C. It is always deterministic and can be predicted with complete accuracy.
- D. It is irrelevant as long as the reward function is well-defined.
- E. It changes over time and does not maintain stationarity.

Question 3

What is the purpose of introducing stochasticity (uncertainty) in actions within a Markov Decision Process?

- A. To simplify the computation of the optimal policy.
- B. To model real-world scenarios where actions do not always lead to predictable outcomes.
- C. To ensure that all actions have the same probability of success.
- D. To eliminate the need for a transition model in decision-making.
- E. To focus solely on immediate rewards rather than long-term outcomes.

Question 4

How is the utility of a sequence of states in an MDP typically calculated?

- A. By counting the number of states visited.
- B. By multiplying the rewards of each state.
- C. By adding up all the rewards received in those states.
- D. By averaging the rewards over the total number of states.
- E. By considering only the rewards of the final state in the sequence.

Question 5

What is the primary goal of using discounted rewards in Markov Decision Processes?

- A. To increase the importance of immediate rewards over future rewards.
- B. To ensure that the sum of rewards over an infinite sequence is finite.
- C. To prioritize actions that lead to higher rewards in the shortest time.
- D. To eliminate the need for a policy in decision-making.
- E. To simplify the calculation of the transition model in an MDP.

Answer Key

1. B
2. B
3. B
4. C
5. B

Answer Key

1. A
2. C
3. B
4. B
5. B